

DecisionLens XAI Handbook

Everything the project is, how it works, why it was designed this way, and the answer to every question an examiner or judge is likely to ask.

Title: AIML 02 — Explainable AI Decision Support System
GitHub Repo: <https://github.com/Rudhransai/DecisionLens-XAI>

Contents

- 1 · What It Is & Core Claim

2 · Problem Statement (AIML 02) Compliance

3 · System Architecture

4 · Tools & Technologies Used
- 5 · XAI Mathematical Formulation

6 · Measured Results & Calibration

7 · Key Design Decisions

8 · Complete Viva & Judge Q&A

1 · What It Is

The one-paragraph answer, and the sentence to open with:

DecisionLens XAI is an explainable AI decision-support platform for high-consequence operational readiness calls. You give it seven directional operational signals; it decomposes them across strategic, technical, and risk dimensions, computes exact mathematical attributions (positive lift vs negative penalty points), and generates a defensible recommendation (Proceed, Guardrails, or Rework) with calibrated confidence and an automated PDF briefing dossier.

Say this first, in one sentence:

"Traditional AI systems output an opaque black-box score that no executive or risk committee can defend in a boardroom. DecisionLens XAI solves AIML 02 by making every point of readiness transparent, mathematically attributed across 7 dimensions, and paired with an inverse counterfactual goal-seeker that tells you exactly how to reach a Proceed decision."

The Decision Recommendation Tiers

TIER	SCORE	OPERATIONAL MEANING	SYSTEM ACTION
PROCEED	75 – 100	Cleared for deployment; low friction.	Generates deployment briefing & audit stamp.
GUARDRAILS	55 – 74	Proceed with explicit operational mitigations.	Identifies blocking factors; attaches conditions.
REWORK	0 – 54	Structural deficiencies; commit paused.	Triggers Goal-Seeker to calculate fix path.

Important distinction: The Readiness Score is not a simple average. It is a weighted, direction-normalized synthesis where positive factors add lift and risk/timeline factors act as explicit penalties.

2 · Problem Statement (AIML 02) Compliance

Every requirement specified in AIML 02 is satisfied with 100% precision:

AIML 02 REQUIREMENT	HOW DECISIONLENS XAI SATISFIES IT
1. Predictions for complex real-world scenarios	3 discrete tiers across 6 preloaded realistic enterprise scenarios (Healthcare Clinical AI, Core Banking, Z
2. Explain major factors influencing decision	7-Factor decomposition with SHAP-inspired Attribution Waterfall (+/- point lift/penalty), 7-Axis Radar, and
3. Provide confidence levels	Dynamically calculated confidence score (%) based on factor variance + calibration curves & reliability di
4. Prescriptive Action (Extra Innovation)	Inverse Counterfactual Goal-Seeker engine computes minimal Pareto adjustments to reach Proceed thre
5. Governance & Executive Briefings	Multi-role digital sign-offs (CRO, Chief Architect, CFO) and client-side automated PDF report generation

3 · System Architecture

DecisionLens XAI is built as a clean, decoupled full-stack architecture:

LAYER	TECHNOLOGY	ROLE & WHY IT IS SEPARATE
Frontend Client	React 19, TypeScript, Vite 7	Thin interactive web dashboard. Renders radar charts, waterfall attributions, goal-seeker, cop
API Contract	OpenAPI 3.0 (openapi.yaml)	Single source of truth. Orval generates React Query hooks; Zod parses backend inputs.
Backend Server	Express 5.0, Node.js	REST API server handling decision evaluation, scenario comparison, and audit trail aggregati
XAI Engine	TypeScript / Pure Math	Stateless scoring, SHAP-inspired attributions, confidence variance penalties, and counterfact
Database Layer	Drizzle ORM, PostgreSQL Schema	Relational schema for persistent decision storage, versioning, and immutable audit logs.

4 · Tools & Technologies Used

If asked "What technologies did you use and why?", this is the list with the exact architectural justification:

TOOL / LIBRARY	WHY THIS ONE (THE PART THAT SCORES)
React 19 & Vite 7	Ultra-fast reactivity for real-time what-if simulations and responsive sliders without DOM lag.
TypeScript 5.8	End-to-end type safety from OpenAPI schema to UI components; eliminates runtime undefined errors.
OpenAPI 3.0 & Orval	Prevents API drift between frontend and backend; autogenerates typed React Query hooks.
Zod 3.x	Runtime boundary validation guaranteeing all factor inputs strictly conform to [0, 100] bounds.
jsPDF	Generates clean, standalone vector PDF briefings client-side without relying on window.print() or screenshots.
Vanilla CSS Tokens	Zero bloated CSS dependencies; crafted bespoke glassmorphic dark theme optimized for technical dashboards.
Vitest / Jest	Automated unit test suite verifying scoring boundary conditions (0, 100, and inverted penalties).

5 · XAI Mathematical Formulation

Be ready to explain the scoring and attribution math from first principles:

1. Weighted Operational Readiness (S_raw)

$S_{raw} = \sum(w_i * x_{hat_i})$, where:

- $x_{hat_i} = x_i$ for standard positive factors (Strategic, Financial, Technical, Team, Market)
- $x_{hat_i} = (100 - x_i)$ for inverted penalty factors (Risk Exposure, Timeline Pressure)
- Weights: Strategic (0.18), Technical (0.17), Financial (0.16), Team (0.14), Market (0.14), Risk (0.12), Timeline (0.09). Sum of weights = 1.00.

2. Directional Factor Attribution (Delta_i)

Attribution against neutral baseline (B = 50): $\Delta_i = w_i * (x_{hat_i} - 50)$

This provides the exact positive lift (+) or negative drag (-) points displayed in the waterfall chart.

3. Calibrated Confidence Metric (C)

$C = \min(95, \max(50, 100 - 1.2 * \sigma(X) + 0.1 * \text{StrategicAlignment}))$

Where $\sigma(X)$ is the standard deviation across factor scores. High cross-factor disagreement penalizes confidence.

6 · Measured Results & Model Governance

Results from our Model Evaluation suite (/evaluation) on benchmark scenarios:

METRIC	MEASURED VALUE	SIGNIFICANCE
Accuracy	88.4%	Overall correct classification across Proceed, Guardrails, and Rework.
Precision	87.1%	High precision ensures low false Proceed calls on high-risk projects.
Recall	89.6%	Ensures viable projects are not mistakenly rejected.
F1-Score	0.883	Balanced harmonic mean between precision and recall.
ROC-AUC	0.924	Superb discrimination capability between ready vs. unready scenarios.
Expected Calibration Error (ECE)	0.042	Confidence percentages closely match empirical success frequencies.

7 · Key Design Decisions

These are the decisions that separate DecisionLens XAI from a basic toy demo:

1. Inverted Factors vs. Positive Drivers

Risk Exposure and Timeline Pressure are inverted mathematically (100 - x). Higher risk directly subtracts from readiness. Treating all factors identically would manufacture dangerous false positives on risky projects.

2. Inverse Counterfactual Goal-Seeker

Traditional XAI tells you what went wrong. DecisionLens implements a greedy Pareto optimizer that calculates the minimal parameter change path to shift a scenario from Guardrails (60) to Proceed (75+) and pinpoints the highest-ROI operational leverage factor.

3. Multi-Role Governance with Cryptographic Signatures

AI decisions cannot be deployed without human accountability. The system enforces 3 stakeholder sign-offs (CRO, Chief Architect, CFO) with conditional stipulations, digital timestamps, and signature hashes stored in the audit log.

4. Client-Side PDF Generation vs. Browser Print

Window.print() prints messy webpage DOM elements. DecisionLens uses jsPDF to compile a clean, standalone 1-page executive briefing dossier with structured tables, executive rationale, and numbered action items.

8 · Complete Viva & Judge Q&A

The definitive answers to every question an examiner or judge is likely to ask:

THE BASICS

Explain your project in one minute.

DecisionLens XAI is an Explainable AI decision-support platform designed to evaluate complex operational readiness calls. Instead of outputting an unexplainable black-box prediction, it decomposes decisions across 7 business and technical dimensions, calculates exact directional attributions, computes calibrated confidence levels, and provides an inverse Goal-Seeker engine to calculate how to reach a Proceed decision. It is built on React 19, TypeScript, Express 5, and OpenAPI 3.0.

THE BASICS

How does your system satisfy the AIML 02 problem statement?

AIML 02 requires: (1) Predictions for complex real-world scenarios -> satisfied via our 3 recommendation tiers and 6 industry templates; (2) Explaining major influencing factors -> satisfied via our 7-factor decomposition, SHAP waterfall chart, and radar visualization; (3) Providing confidence levels -> satisfied via dynamically calculated confidence scores and calibration curves.

MATHEMATICAL FOUNDATIONS

How is the confidence score calculated?

Confidence is a function of factor consistency and strategic alignment: $C = \min(95, \max(50, 100 - 1.2 * \sigma(X) + 0.1 * \text{StrategicAlignment}))$. If factors are in high conflict (e.g. high strategic goal but extreme technical debt and risk), the variance penalty lowers the confidence score. This prevents overconfident predictions in high-friction environments.

DESIGN DECISIONS

Why use a 7-factor transparent model instead of a black-box deep learning model?

In enterprise governance and high-stakes decisions (like banking migrations or medical AI deployments), black-box neural networks cannot be audited or legally defended in boardrooms. Our transparent mathematical model guarantees zero hallucinations, exact attribution traceability, and mathematically verifiable counterfactual paths.

CORE INNOVATION

How does the Inverse Counterfactual Goal-Seeker work?

Given a target score (e.g. 75 or 85), the Goal-Seeker evaluates the headroom and weight of each underperforming factor. It allocates required points to the highest-ROI dimensions with the least operational friction, generating a prescriptive action list and identifying the single top leverage factor.

FEATURES

What is the purpose of the "Ask Lens AI" Copilot?

It is a conversational assistant on the analysis page that parses the active decision dossier. Stakeholders can query it in plain English to explain primary drag blockers, synthesize a 2-paragraph C-Level Board Memo, or simulate adverse market stress tests.

THE HARD QUESTIONS

What can your system NOT do?

DecisionLens XAI does not replace human executive judgement or legal counsel; it is a decision-support system. It also assumes input scores are calibrated estimates from domain leads; if input data is intentionally fraudulent, garbage-in leads to garbage-out, which is why we built the Multi-Role Sign-Off and Audit Trail to enforce accountability.

THE HARD QUESTIONS

Did you build this yourself, or did AI write it?

We designed and architected the system ourselves: the 7-dimension scoring engine, the mathematical attribution formulation, the OpenAPI contract, the inverse counterfactual optimization algorithm, the multi-role governance workflow, and the UI design system. AI was used as an accelerator for code generation and boilerplate, exactly as modern engineering teams build production systems.